

# Assignment 2: From Trees to Boosting Power

**Due Date:** Wednesday, October 22, 2025 at 12:00 AM (midnight, Eastern Time)

**Policy:** <https://columbia-coms4995.github.io/aml-fall2025/#policies>

## Submission

- A runnable Jupyter/Colab Notebook (`.ipynb`)
- A PDF report (5-7 pages with clear structure and visualizations)

## Objective

This assignment expands upon A1 and challenges you to explore the **power of Boosting methods**. You will move beyond a single model and investigate how Boosting approaches improve generalization by correcting residual errors in sequential learners.

By the end of this assignment, you should be able to:

- Train and tune tree-based and boosting models.
- Analyze the bias-variance trade-off through validation results and boosting dynamics.
- Visualize feature importance and interpret how boosting models evolve during training.

## Datasets

Choose one of the following datasets:

### Bank Marketing Dataset

**Task:** Predict whether a customer will subscribe to a term deposit based on demographic and campaign features.

**Link:** <https://archive.ics.uci.edu/dataset/222/bank+marketing>

### Home Credit Default Risk Dataset

**Task:** Predict whether an applicant will repay a loan, using a wide range of application, credit, and behavioral data.

**Link:** <https://www.kaggle.com/c/home-credit-default-risk>

Both datasets are suitable for tree-based and boosting algorithms, featuring mixed data types, non-linear relationships, and moderate feature dimensionality.

## Steps

### 1. Data Preparation

- a. Handle missing values and inconsistent types.
- b. Apply label encoding or one-hot encoding to categorical variables.
- c. Optionally construct new features (e.g., ratios, age groups, spending levels).
- d. Split data into train / validation / test sets (e.g., 70 / 15 / 15).

### 2. Baseline Model: Decision Tree

- a. Train a `DecisionTreeClassifier` using Gini impurity.
- b. Tune `max_depth`, `min_samples_split`, and `ccp_alpha` (pre-pruning).
- c. Evaluate train and test accuracy to illustrate overfitting vs underfitting.
- d. Plot the learning curve to show how training and validation scores change with depth.

### 3. Boosting Methods

Your main focus in this assignment is to **understand and implement boosting algorithms**, and compare them with a baseline tree-based method.

#### a. Baseline

Train either a simple **Decision Tree** or a **Random Forest** as your baseline model. Tune its depth, minimum samples, and regularization parameters. Record training and validation accuracy.

#### b. Gradient Boosting

Train a **GradientBoostingClassifier** using scikit-learn. Explore how the following parameters affect performance and overfitting:

- `learning_rate`
- `n_estimators`
- `max_depth`
- `subsample`

#### c. XGBoost

Install and use **XGBoostClassifier** (`xgboost` library). Perform similar tuning and visualize:

- training vs. validation loss (using `eval_set` or early stopping)
- feature importance (`plot_importance`)
- effect of learning rate and tree depth on bias-variance balance

#### d. Optional (Bonus)

Implement [LightGBM](#) and compare their speed and accuracy against XGBoost. Discuss how different boosting frameworks handle categorical features and regularization.

### 4. Model Evaluation and Visualization

a. Metrics: Accuracy, Precision, Recall, F1, AUCPR.

b. Plot:

- Confusion matrix
- PR curve
- Feature importance (bar chart)
- Learning curve (bias-variance illustration)

c. Discuss differences between baseline Decision Tree vs Boosting performance.

Additionally, for your **boosting models**, plot and analyze:

- **Training vs Validation Loss over Boosting Iterations** (e.g., using `eval_set` or custom matplotlib visualization)
- **Effect of Learning Rate:** Compare results for 0.01, 0.1, and 0.3
- **Bias-Variance Trade-off:** Discuss how increasing the number of estimators affects generalization

These visualizations will help you interpret how boosting methods incrementally reduce bias while controlling variance.

### 5. Discussion and Interpretation

- Explain why boosting methods improve generalization.
- Relate your results to the bias-variance trade-off.
- Comment on computational cost and interpretability.
- Reflect on how feature importance helps interpret predictions.

### 6. Report Writing (6 - 10 pages)

- Introduction:** Problem definition & dataset description.
- Methods:** Preprocessing and modeling steps (summary table optional).
- Results:** Evaluation metrics and visualizations.
- Discussion:** Interpretation of findings, limitations, bias-variance reflection.
- AI Tool Disclosure:** List any AI tools (e.g., ChatGPT, Gemini, Claude) and their roles.

### 🔑 Pointers & Hints

- Start with `DecisionTreeClassifier` or `RandomForestClassifier` as a baseline before moving to boosting.
- Use `GridSearchCV` or `RandomizedSearchCV` for hyperparameter tuning.

- Visualize feature importance using `model.feature_importances_` or `xgb.plot_importance()`.
- `cross_val_score`, `learning_curve`, `validation_curve` in scikit-learn are helpful.
- For XGBoost, use the `eval_set` and `early_stopping_rounds` parameters to monitor model convergence and prevent overfitting.
- Control random seeds for reproducibility.

## Grading Rubric

Category	Weight	Description
Data Preparation & Baseline Decision Tree	15	Proper preprocessing and baseline model tuning.
<b>Boosting Implementation (Gradient Boosting / XGBoost)</b>	<b>30</b>	Correct setup, parameter tuning, and visualization of training dynamics.
Evaluation & Visualization	20	Metrics (Accuracy, F1, PR), plus boosting loss curves and feature importance.
Discussion & Interpretation	20	Strong explanation of how boosting reduces bias and improves generalization.
AI Tool Usage Disclosure	10	Transparent acknowledgment of AI tools and your personal contributions.
Bonus (+ up to 10%)	-	Compare XGBoost vs LightGBM; analyze effect of learning rate and number of estimators.