

COMS4995W32

Applied Machine Learning

Dr. Spencer W. Luo

Columbia University | Fall 2025



Generative vs. Discriminative Approaches: Naive Bayes vs. Linear Regression

Agenda



- Motivation
- Naive Bayes (Generative)
- Linear Regression (Discriminative)
- Case Study



G vs. D: The Big Picture

- Generative approach (e.g., Naive Bayes)
 - Learns how the data is generated
 - Like: “learning the recipe 🍳”
- Discriminative approach (e.g., Linear Regression)
 - Directly learns the mapping
 - Like: “tasting and giving a score 🍴🍽️”
- Same goal → **make predictions**
 - Yet 2 very different philosophies 📖

Generative Philosophy





Learns the story of how data is produced (**joint distribution**)

- Examples:
 - Naive Bayes
 - assumes words are independent
 - builds probability recipe for each class
 - Language models
 - predict next word by modeling sentence generation



Generative Philosophy

✓ Strengths of Naive Bayes

-  Performs well on small datasets (low variance, simple structure)
-  Strong baseline - often hard to beat with simple methods

✗ Limitations of Naive Bayes

-  Naive assumption: features are conditionally independent
-  Fails if features are correlated (e.g., “New” + “York”)

Discriminative Philosophy





Learns the **boundary** or function that best predicts outcomes

- Examples:
 - Linear Regression (fit a line directly to minimize error)
 - Logistic Regression, SVM, modern neural networks





Discriminative Philosophy

Strengths of Discriminative Models

-  Directly optimize prediction accuracy - learn $p(y|x)$
-  Often higher accuracy than generative models (fewer assumptions)

Limitations of Discriminative Models

-  Cannot model data generation $p(x,y)$
-  Need more data than generative models to perform well

Same Goal, Different Roads.



Why This Matters?




Different learning philosophies → different trade-offs

- Generative = Once we need to understand data generation
- Discriminative = When prediction accuracy is the only goal


Discriminative vs. Generative

Discriminative



- Only care about estimating the conditional probabilities
- Very good when underlying distribution of data is really complicated (e.g. texts, images, movies)

Generative



- Model observations (x,y) first, then infer $p(y|x)$
- Good for missing variables, better diagnostics
- Easy to add prior knowledge about data



Naive Bayes (Generative)



Bayes' Rule

- Prediction is based on updating belief/class after seeing evidence

$$p(y|x) = \frac{p(x|y)p(y)}{p(x)}$$




- **Prior** $p(y)$: how common each class is
- **Likelihood** $p(x|y)$: how features look under each class
- **Posterior** $p(y|x)$: updated belief after seeing
- 💡 belief updating → start with a guess $p(y)$, refine with data $p(x|y)$



Naive Bayes Assumption

- Assuming features are **conditionally independent** given the class

$$p(x|y) = \prod_{j=1}^d p(x_j|y)$$

- Big simplification: turns complex joint distribution into a product
- Works surprisingly well for various applications
 -  each pixel \approx independent feature
 -  word presence treated as independent
- Intuition: Given the class, every word casts its own vote 



Example

Suppose we want to **classify** emails as **Spam** ($y=1$) or **Ham** ($y=0$).

Vocabulary (features): {lottery, meeting, beef}

Thus each feature x_j stands for whether that word appears in the email (0/1)

👎 “Win a free lottery ticket” $\rightarrow x = [1, 0, 0]$


👍 “We will have a meeting” $\rightarrow x = [0, 1, 0]$

👍 “beef sometimes holds a meeting” $\rightarrow x = [0, 1, 1]$

🤔 “You have won a lottery!”



Bernoulli Naive Bayes (binary features)

- Features: binary word presence/absence (0 or 1)
- Each word = yes/no vote 
 - If present → contributes evidence
 - If absent → neutral or negative evidence
- Common for short text (SMS, tweets, headlines)
- Formula

$$P(x_j \mid y) = \theta_{j|y}^{x_j} (1 - \theta_{j|y})^{1-x_j}$$



Learning Parameters (MLE)

- Class prior estimated by class frequencies:

$$\hat{P}(y) = \frac{N_y}{N}$$

- Word probabilities estimated by frequency inside each class:

$$\hat{\theta}_{j|y} = P(x_j=1|y) = \frac{N_{j,y}}{N_y} = \frac{\text{Count of word } x_j \text{ in class } y}{\text{Total word counts in class } y}$$

Bernoulli Naive Bayes (binary features)



$$P(x_j = 1 \mid y) = \hat{\theta}_{j|y} \quad P(x_j = 0 \mid y) = 1 - \hat{\theta}_{j|y}$$

$$P(x_j \mid y) = \theta_{j|y}^{x_j} (1 - \theta_{j|y})^{1-x_j}$$



Smoothing (Laplace / Add- α)

Problem: unseen word \rightarrow zero probability

Solution: add α to counts

$$\hat{\theta}_{j|y} = \frac{N_{j,y} + \alpha}{N_y + \alpha d}$$

where d = vocabulary size

Effect: prevents zeros, makes model more robust



Numerical Stability



Multiplying many small probabilities → **underflow**

- Example: $0.001^{100} \approx 0$ on a computer

Solution: switch to log-space (turn products into sums)

$$\log P(y \mid x) \propto \log P(y) + \sum_{j=1}^d \log P(x_j \mid y)$$

Benefits:

- Avoids numerical underflow
- Computation becomes addition (faster, more stable)

PS: All modern NB implementations use log probabilities internally

Gaussian Naive Bayes (continuous features)



- Features: continuous values (e.g. numeric features)
- Assumption: each feature follows a Gaussian distribution within each class

$$P(x_j|y) = \frac{1}{\sqrt{2\pi\sigma_{j,y}^2}} \exp\left(-\frac{(x_j - \mu_{j,y})^2}{2\sigma_{j,y}^2}\right)$$



Practical Pipeline



Step 1 - Preprocessing

- Tokenize, lowercase, remove stopwords
- Handle missing values if needed



Step 2 - Feature Extraction

- Bag-of-Words / binary counts → Bernoulli NB
- Word counts / TF-IDF → Multinomial NB
- Continuous features → Gaussian NB



Practical Pipeline



Step 3 - Classifier

- Choose model: BernoulliNB, MultinomialNB, GaussianNB
- Learn parameters by frequency counts (MLE + smoothing)



Step 4 - Evaluation

- Metrics: Accuracy, Precision, Recall, Confusion Matrix
- Use `cross_val_score` for validation



When NB Works / Fails



Works well:

- High-dimensional sparse features (text, word counts)
- Small datasets (few samples per class)



Fails when:

- Strong feature correlations
- Complex dependencies between features



Linear Regression



From Generative to Discriminative

- Naive Bayes: model joint distribution $P(x,y)$
- Linear Regression: directly fit $f(x) \approx y$

Do not learn the recipe, just draw the line that fits the taste 

Goal: **minimize classification error**, not model data generation



Hypothesis Form

Linear model = weighted sum of features

$$\hat{y} = \theta^\top x = \sum_{j=1}^d \theta_j x_j + b$$

Interpretation

- Each feature contributes **proportionally** to the prediction
- The weight tells us how much that feature **matters**

Trick

- add $x_0 = 1 \rightarrow$ bias term included
- Ensures model can fit data that does **not** pass through the origin



Loss Function (MSE)

Residual sum of squares (RSS):

$$L(\theta) = \sum_{i=1}^n (y_i - \hat{y}_i)^2$$

Mean Squared Error (MSE):

$$\text{MSE} = \frac{1}{n} \sum_{i=1}^n (y_i - \theta^\top x_i)^2$$

Intuition: penalizes more on **large errors**



Gradient Descent (Iterative Optimization)

Update rule


$$\theta \leftarrow \theta - \eta \nabla_{\theta} L(\theta)$$

where η = learning rate/step size

Works for large datasets & online learning

Cornerstone of DL 

Gradient Descent Intuition





Ordinary Least Squares (Closed Form)

Analytic solution when $X^T X$ invertible:

$$\theta^* = (X^T X)^{-1} X^T y$$

Analytic solution = one-shot computation

Fast for small/medium datasets

Limitation: unstable with collinearity or ill-conditioned



Regularization Motivation

Problems:

- Overfitting with too many features
 - Multicollinearity → unstable coefficients
-
- Solution: add **penalty** to shrink coefficients
-
- Leads to **Ridge** & **Lasso regression**



Ridge Regression

Loss with L-2 penalty:

$$L(\theta) = \frac{1}{n} \sum_{i=1}^n (y_i - \theta^\top x_i)^2 + \alpha \|\theta\|_2^2$$

- Always unique solution
- Coefficients shrink towards zero (but not exactly zero)
- Good when many small/medium effects are expected



Lasso Regression

Loss with L-1 penalty:



$$L(\theta) = \frac{1}{n} \sum_{i=1}^n (y_i - \theta^\top x_i)^2 + \alpha \|\theta\|_1$$

- Encourages sparsity \rightarrow some coefficients exactly 0
- Automatic feature selection
- Harder optimization (non-differentiable at 0)

Ridge vs Lasso (Geometry)



Lasso	Ridge
diamond constraint \rightarrow corners	circular constraint \rightarrow smooth shrinkage



Summary of Linear Regression



- Strengths: simple, interpretable, fast baseline
- Limitations: sensitive to outliers, assumes linearity
- Extensions: logistic regression, neural networks



NB vs. LR

Comparison



Aspect	Naive Bayes (Generative)	Linear Regression (Discriminative)
Model	Learns joint $P(x,y)$	Learns conditional $P(y x)$ directly
Philosophy	Recipe of data → simulate how data is generated	Boundary drawing → just separate outcomes
Features	Like word presence (independent assumption)	Fits weighted sum of features
Strengths	Simple, fast, good for text, small data	Interpretable, flexible, accurate with enough data
Limitations	Independence assumption unrealistic	Sensitive to outliers, needs more data



Case Study