

# **COMS4995W32**

# **Applied Machine Learning**

Dr. Spencer W. Luo

Columbia University | Fall 2025



# Unsupervised Learning

# Machine Learning Paradigm



**Supervised:** learn  $f$  from  $(x, y) \rightarrow$  model **learns mapping**

**Unsupervised:** only  $x \rightarrow$  **discover structure**, or latent features

- It discovers hidden patterns or latent structures in data
- When new data arrives, the model can represent it based on learned structure

**Semi-supervised:** **hybrid** of both worlds

- Uses small labeled set to guide large unlabeled pool.



# Clustering with K-Means

# What is Clustering? 🧒🧒





Goal: group similar data points into clusters **based on similarity**

Each cluster represents a natural “pattern” or “category”

Scenarios: segmenting customers, species, or document topics

🤔 No labels - the algorithm must decide grouping itself

# K Clusters



# Similarity Metrics



Similarity defines **how close** two samples are


Common distances:

- **Euclidean** → geometric closeness
- **Manhattan** → sum of absolute differences
- **Cosine** → angle between vectors (useful for text)


# Comparisons




**KMeans (Euclidean) — clusters by distance (radius)**



**K-Medians (Manhattan/L1) — diamond-like regions**



**KMeans (Cosine) — clusters by direction (angle)**





# The K-Means Clustering Objective



Partition data into K clusters that minimize within-cluster distance

- Each cluster is represented by its **centroid** (mean vector)
- During optimization, each point is assigned to the nearest centroid, and the centroid is recomputed as **the mean of all points in that cluster**:

$$\min_{\{\mu_k\}_{k=1}^K} \sum_{i=1}^N \|x_i - \mu_{c_i}\|^2 \quad \text{where} \quad \mu_k = \frac{1}{|C_k|} \sum_{x_i \in C_k} x_i$$

# Algorithm Overview



Step 1 Initialize K random centroids


Step 2 Assign each point to its nearest centroid

Step 3 Re-calculate centroids as cluster means

Step 4 Repeat until assignments no longer change

Result → clusters + centroids

# Algorithm Overview



# Iterative Refinement



K-Means alternates between two simple steps:

- Assignment: label each point  $\rightarrow$  nearest centroid
- Update: recalc centroid  $\rightarrow$  mean of assigned points

Converges when centroids stop moving

Stops after a few iterations

# K-means Strengths & Limitations



- ✓ Simple and fast – scales well for large data
- ✓ Intuitive – easy to interpret centroids
- ✗ Sensitive to initialization and outliers
- ✗ Assumes spherical clusters, fails for complex shapes
- ✗ Requires manual K selection

# Beyond K-Means



Hierarchical Clustering → tree-like structure

- Builds clusters step by step (divisive), visualized as a hierarchy

DBSCAN → density-based, detects arbitrary shapes

- Groups nearby points with high density and marks outliers as noise

Gaussian Mixture Models → probabilistic soft clustering

- Assigns each sample a probability of belonging to each cluster



# Neural Network Fundamentals

# Why Neural Networks?



Many real-world problems are non-linear → linear models fail

- Relationships in data are curved, complex, and high-dimensional

Neural networks can learn complex functions via composition

- By stacking neurons, we can approximate any function ([Universal Approximation Theorem](#))

Inspired by the brain: neurons connect and adapt weights


- Each neuron adjusts its connection strength, gradually improving the generalizability




# Linear vs Nonlinear



**Linear Problem**



**Nonlinear Problem**



# The Big Picture of Neural Networks



Start from linear models → limited expressiveness

Add **nonlinear activations** → learn complex functions

Stack layers → Multilayer Perceptron (MLP)

Train by **forward** + **backward** passes to minimize **loss**

Tune learning with optimization and regularization

Forms the foundation for CNNs, Transformers and ChatGPTs

# The Big Picture of Neural Networks



Start from linear models → limited expressiveness

Add **nonlinear activations** → learn complex functions

Stack layers → Multilayer Perceptron (MLP)

Train by **forward** + **backward** passes to minimize **loss**

Tune learning with optimization and regularization

Forms the foundation for CNNs, Transformers and ChatGPTs

# The Neuron



## Each neuron

- A neuron receives various inputs and computes a **weighted sum** of them
- This sum is then passed through an **activation** function, which introduces **non-linearity**
- The result of this activation becomes the **output** of the neuron, serving as **input** for the next layer


# The Neuron



Nonlinear activations = flexible decision surfaces

- Without nonlinearity, the entire network would behave like a single linear model
- Nonlinear activations allow the network to learn **complex, curved decision boundaries**, making it capable of modeling intricate relationships in data

# The Neuron



$$y = \sigma\left(\sum_{i=1}^n w_i x_i + b\right)$$

# Activation Functions




**Sigmoid** → squashes outputs (good for probabilities)

**Tanh** → zero-centered, faster convergence

**ReLU** → sparse activation, efficient & dominant today

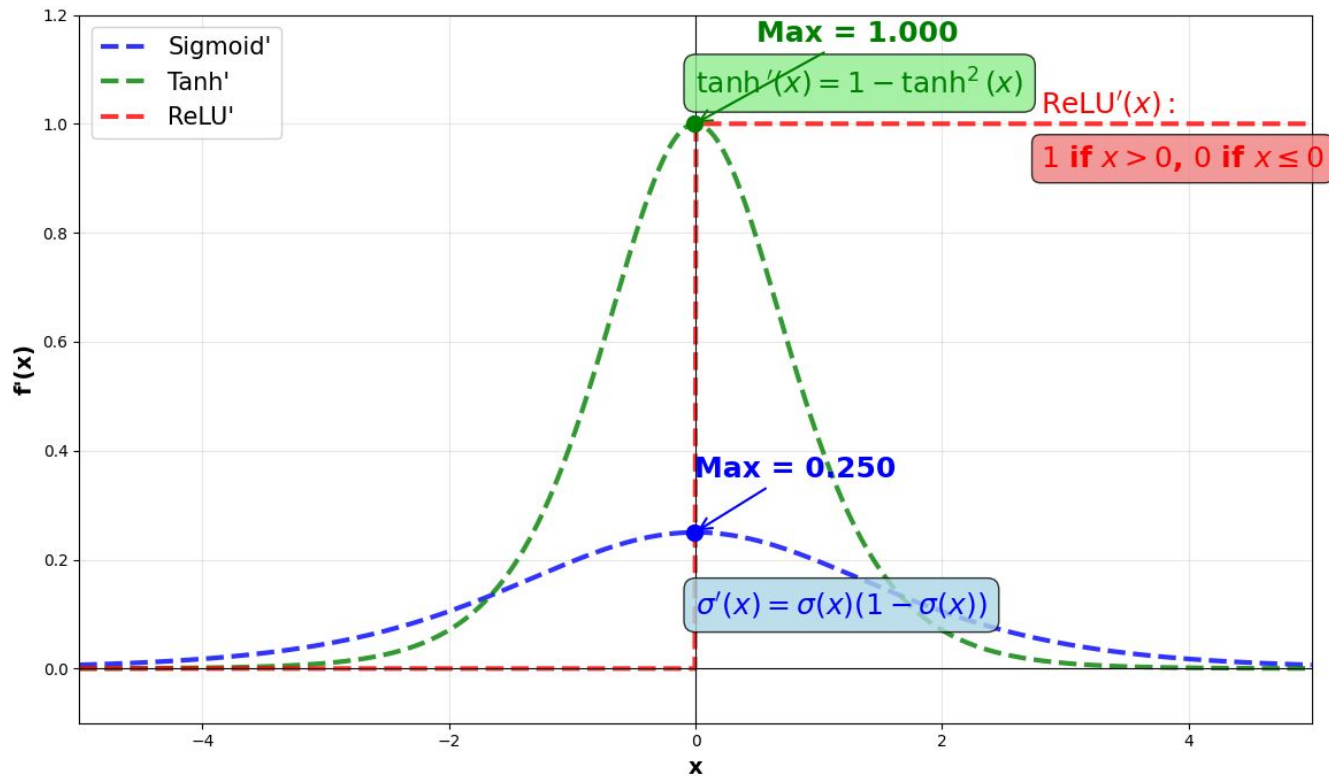
Modern variants: Leaky ReLU, GELU (used in Transformers)

# Activation Functions Comparison





# Derivatives Comparison



# Multilayer Perceptron (MLP)



## Stack of Neurons

- A neural network is built by stacking many simple neurons (linear + activation) into layers

## Hidden Layers → Feature Learning

- Each hidden layer transforms the data into a new representation - from raw input features to increasingly abstract ones

# Multilayer Perceptron (Feed-forward)



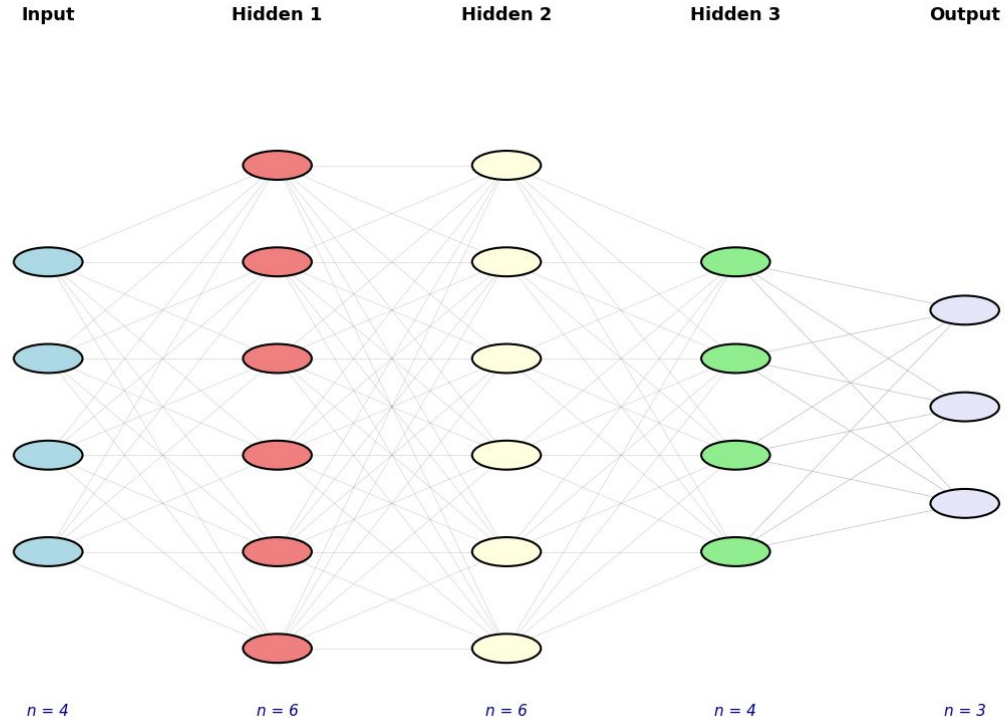
## Output Layer → Task-Specific Prediction

- The final layer converts the learned representation into task-relevant outputs:
  - Regression → real values (e.g., price, temperature)
  - Classification → probabilities via Softmax (e.g., cat / dog / car)

 As we go deeper, the network is not simply memorizing

- It **automatically** represents data using more **abstract and meaningful features**

# Feedforward Architecture



# Training Feedforward Network



Feedforward (left  $\rightarrow$  right)

- Compute layer activations

Loss function (Distance between prediction and ground-truth)

- Map distance to a scalar error

Backprop (right  $\rightarrow$  left)

- Use the computational graph and chain rule to compute gradients

# Training Feedforward Network



Feedforward (left  $\rightarrow$  right)


- Compute layer activations

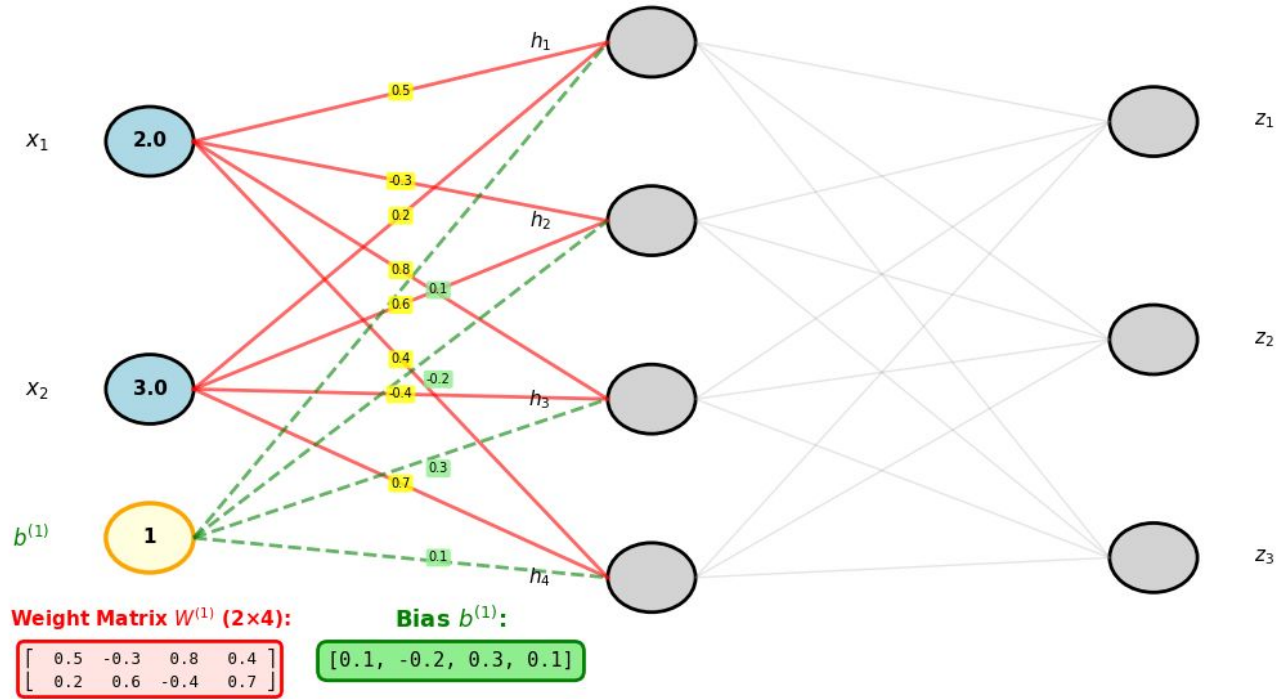
Loss function (Distance between prediction and ground-truth)

- Map distance to a scalar error

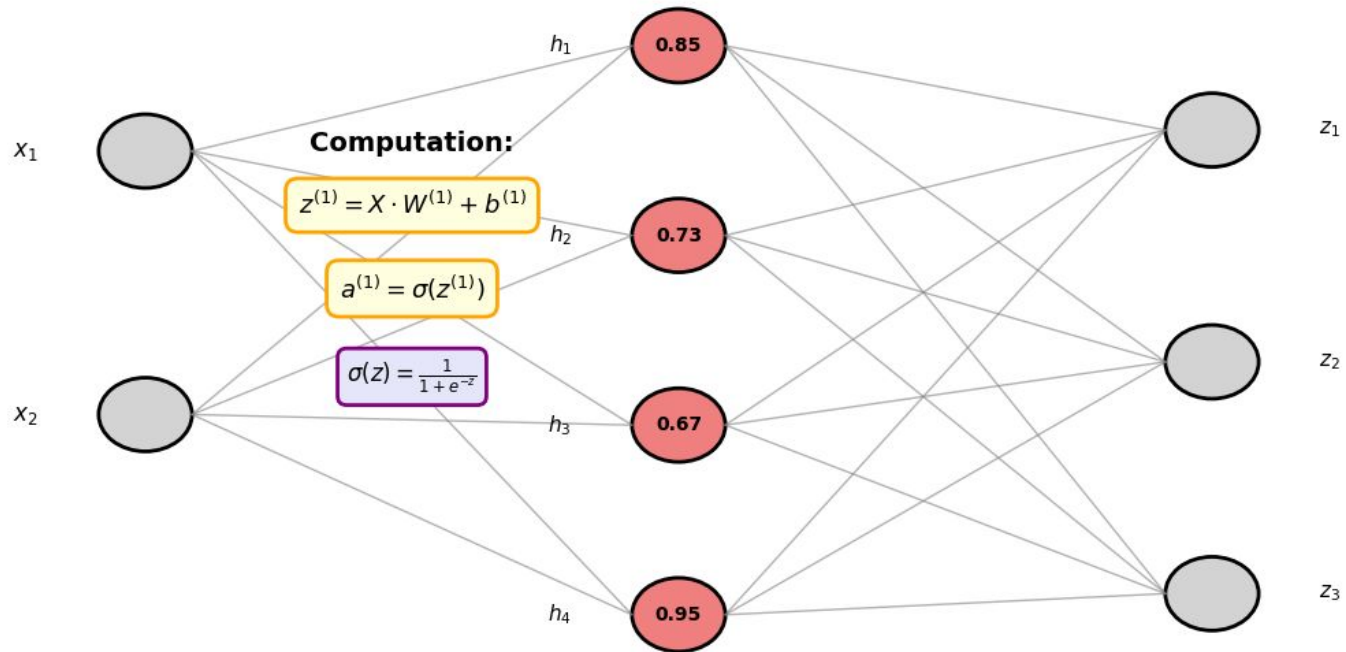
Backprop (right  $\rightarrow$  left)

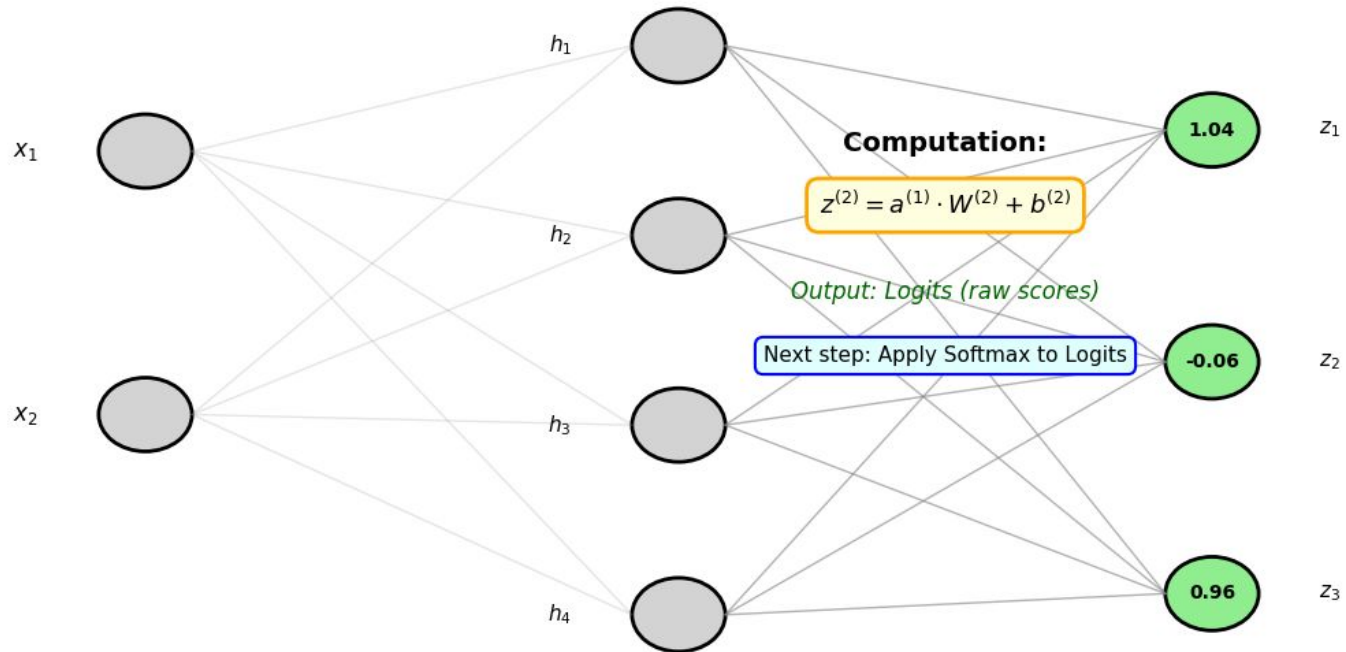
- Use the computational graph and chain rule to compute gradients











# Training Feedforward Network



Feedforward (left  $\rightarrow$  right)

- Compute layer activations

Loss function (Distance between prediction and ground-truth)

- Map distance to a scalar error

Backprop (right  $\rightarrow$  left)

- Use the computational graph and chain rule to compute gradients

# Loss Functions



Quantify how far predictions are from targets

Regression  $\rightarrow$  MSE 
$$\text{MSE} = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2$$

Classification  $\rightarrow$  Cross-Entropy 
$$\text{CrossEntropy} = - \sum_{i=1}^n y_i \log(\hat{y}_i)$$

Lower loss = better model fit on training data

### Input: Logits

$z_1 = 1.04$

$z_2 = -0.06$

$z_3 = 0.96$

### Softmax Function:

$$\text{Softmax}(z_i) = \frac{e^{z_i}}{\sum_{j=1}^K e^{z_j}}$$

*Converts logits to probabilities*

### Properties:

• All probabilities:  $0 \leq p_i \leq 1$

• Sum to 1:  $\sum_i p_i = 1$


### Output: Probabilities

$p_1 = 0.443$

$p_2 = 0.148$

$p_3 = 0.409$

### Probability Distribution:



### Predicted Probabilities

$$\hat{y}_1 = 0.100$$

$$\hat{y}_2 = 0.700$$

$$\hat{y}_3 = 0.200$$

### True Labels (one-hot)

$$y_1 = 0$$

$$y_2 = 1$$

$$y_3 = 0$$

← Target

### Loss Contribution

$$0 \times \log(0.100)$$

$$= -0.000$$

$$-1 \times \log(0.700)$$

$$= 0.357$$

$$0 \times \log(0.200)$$

$$= -0.000$$

$$L = - \sum_{i=1}^K y_i \log(\hat{y}_i)$$

**Total Loss:**  $L = 0.3567$

*Only the correct class contributes to loss*

# Training Feedforward Network



Feedforward (left  $\rightarrow$  right)

- Compute layer activations

Loss function (Distance between prediction and ground-truth)

- Map distance to a scalar error

Backprop (right  $\rightarrow$  left)

- Use the computational graph and chain rule to compute gradients



# Recap Gradient Descent

Update rule

$$\theta \leftarrow \theta - \eta \nabla_{\theta} L(\theta)$$

where  $\eta$  = learning rate/step size

Works for large datasets & online learning

Cornerstone of DL 



# Backpropagation



Efficient gradient computation using chain rule

- Backpropagation applies the chain rule to compute derivatives layer by layer

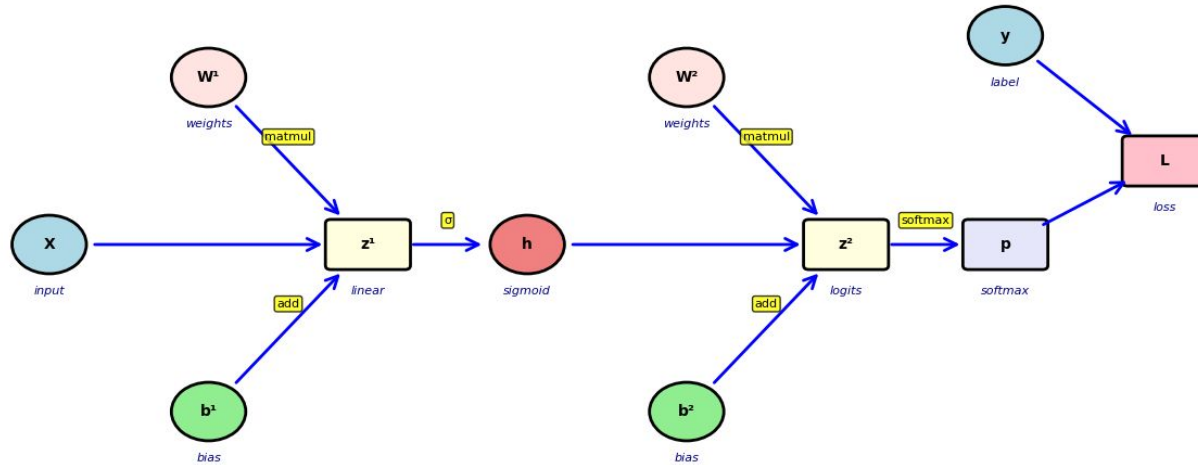
Gradients flow backward from output to input

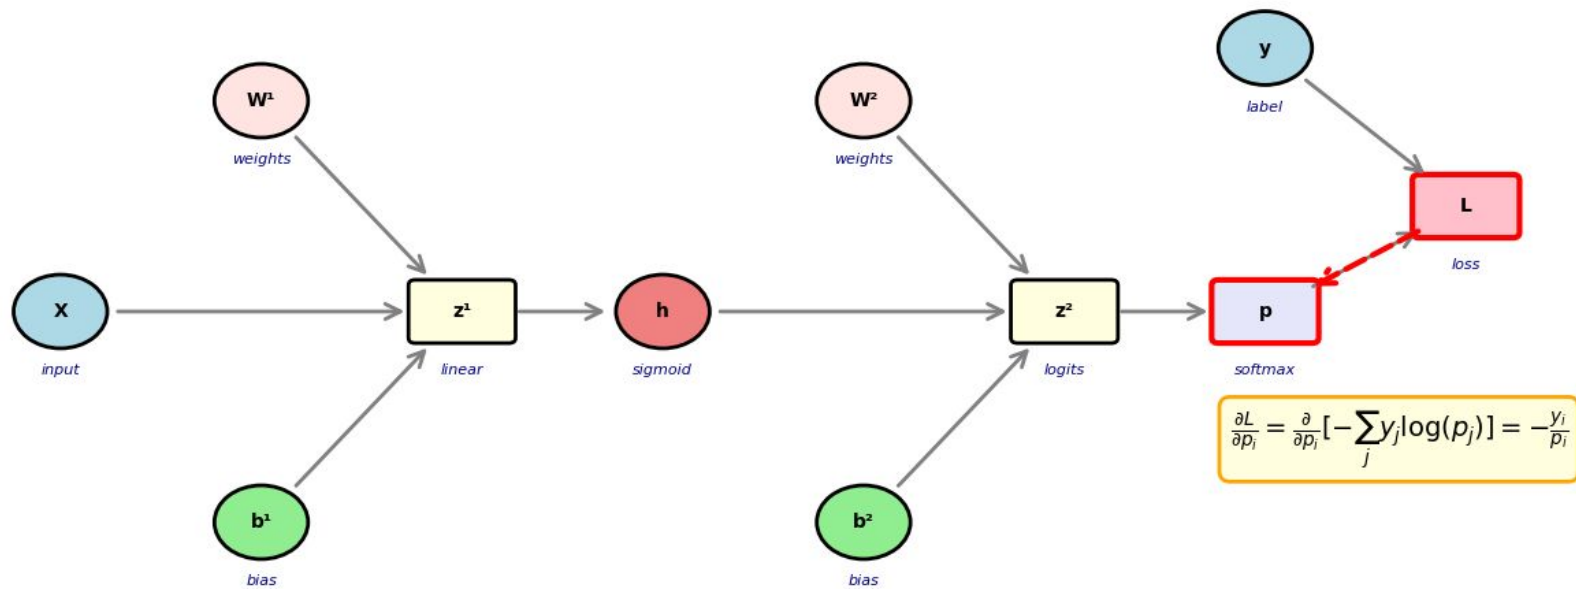
- Each layer receives gradients from the next layer to update its own weights


Enables training of deep networks with millions of weights

- Makes large-scale optimization computationally feasible
- Without backpropagation, deep learning would be impossible to train efficiently

# Computational Graph







$$\frac{\partial L}{\partial z_i^{(2)}} = \sum_j \frac{\partial L}{\partial p_j} \cdot \frac{\partial p_j}{\partial z_i^{(2)}}$$

$$= \sum_j \left( -\frac{y_j}{p_i} \right) \cdot \frac{\partial p_j}{\partial z_i^{(2)}}$$

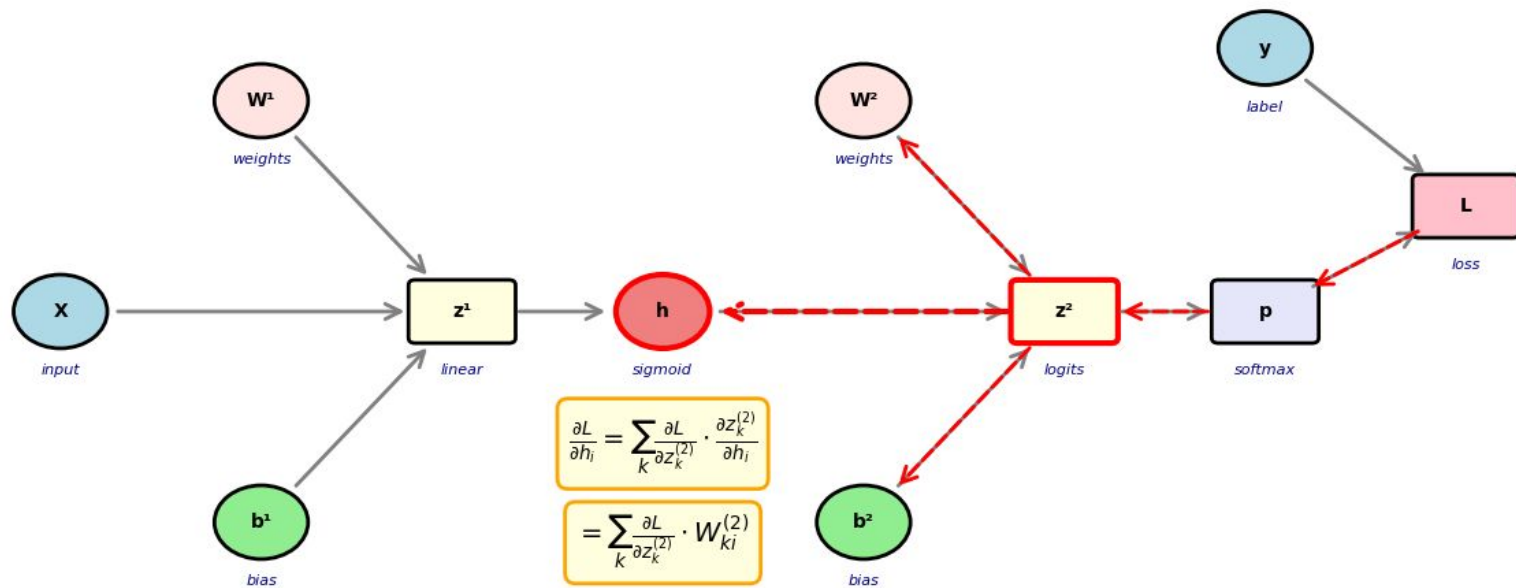
**Softmax derivative:**

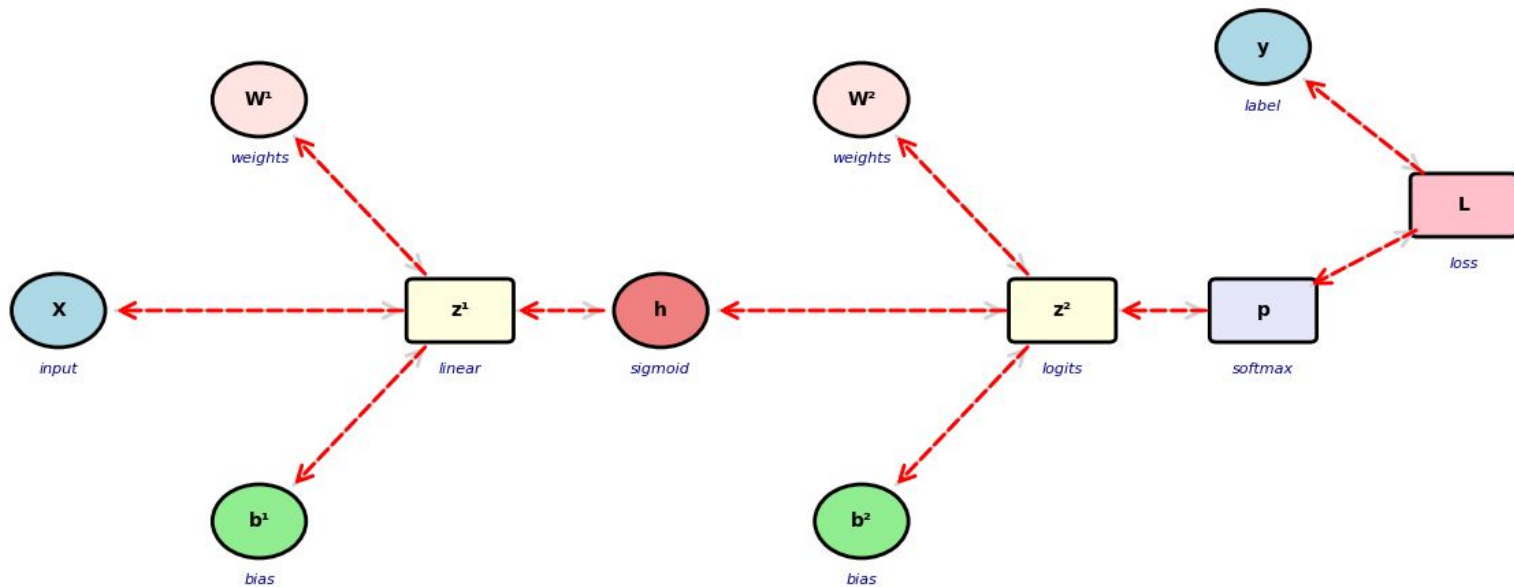
$$\text{if } i = j: \frac{\partial p_i}{\partial z_i} = p_i(1 - p_i)$$

$$\text{if } i \neq j: \frac{\partial p_j}{\partial z_i} = -p_i p_j$$

**Final result:**

$$\frac{\partial L}{\partial z_i^{(2)}} = p_i - y_i$$





### Backpropagation Complete!

✓ Backward Pass: Compute all gradients (red dashed)

✓ Gradients computed via chain rule recursively

✓ Update:  $\theta \leftarrow \theta - \eta \cdot \partial L / \partial \theta$  (Gradient Descent)

# Optimizers



Stochastic GD  $\theta \leftarrow \theta - \eta \nabla_{\theta} L(\theta)$

Too small  $\rightarrow$  slow  too large  $\rightarrow$  divergence 

**Momentum**  $\rightarrow$  Accelerates learning by smoothing gradients with past updates

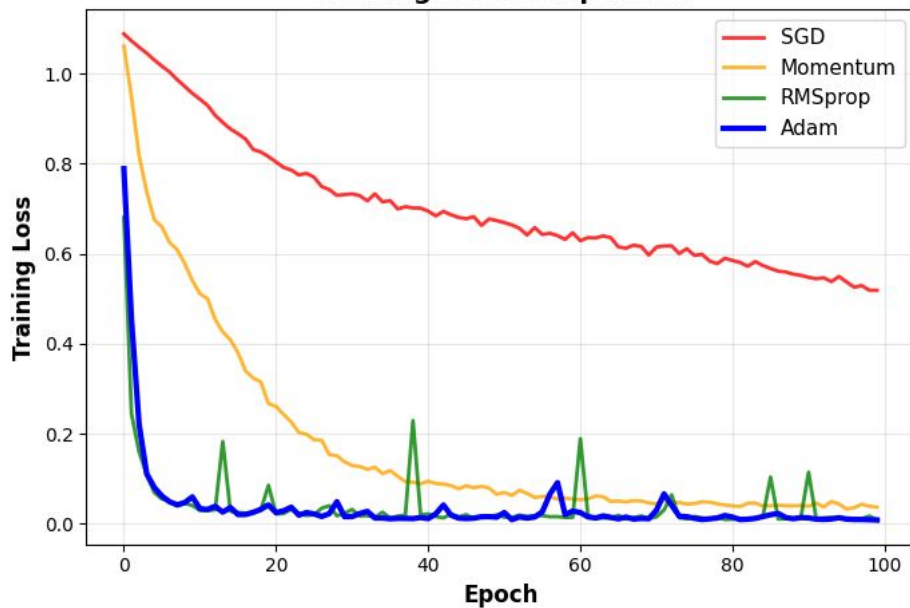
**RMSProp**  $\rightarrow$  Adapts learning rates based on recent gradient magnitudes

**Adam**  $\rightarrow$  Momentum + RMSProp

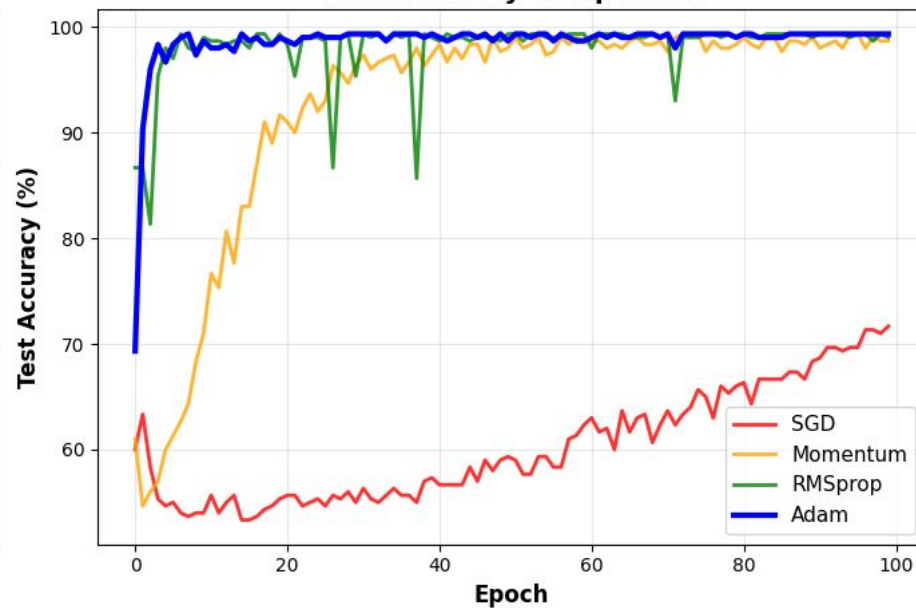
# Comparisons



Training Loss Comparison



Test Accuracy Comparison





# Recap



Start from linear models → limited expressiveness

Add **nonlinear activations** → learn complex functions

Stack layers → Multilayer Perceptron (**MLP**)

Train by **forward** + **backward** passes to minimize **loss**

Tune learning with optimization and regularization

Forms the foundation for CNNs, Transformers and ChatGPTs