

COMS4995W32

Applied Machine Learning

Dr. Spencer W. Luo

Columbia University | Fall 2025



LLM Pre-training & Fine-tuning



Agenda

- Architecture recap
- LLM Pre-training
- LLM Fine-tuning
- Industry trends



Encoder vs Decoder



Decoder-only Family – Language Model (LM)



Mask: **Causal** → token t can observe only tokens $< t$

Objective: **Predict next token** → $P(\text{token} \square | \text{tokens}_{(1\dots\square-1)})$

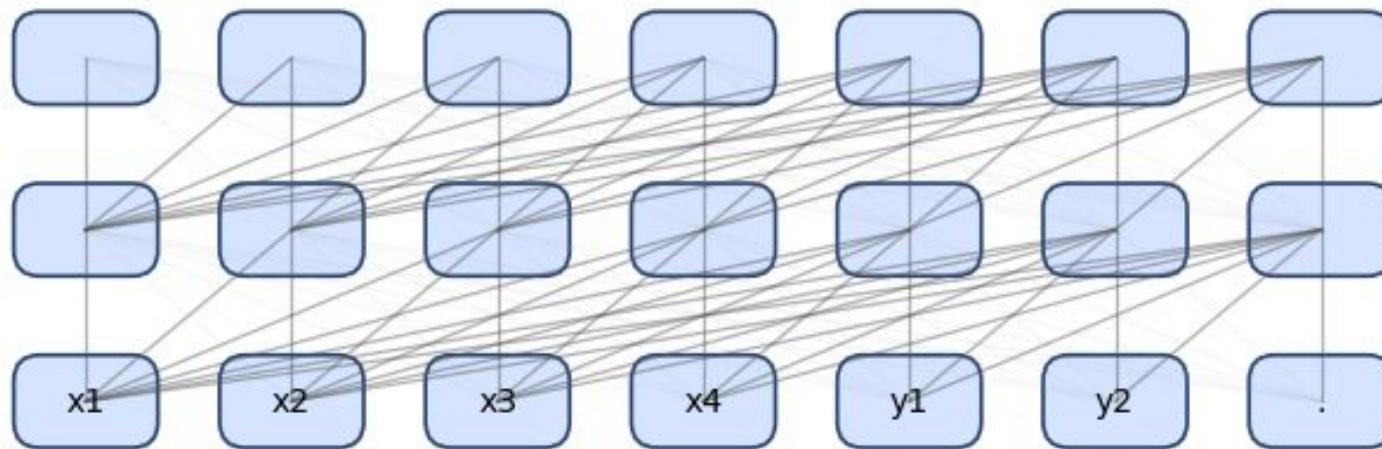
Key: “Generates” but does **not** encode full context

Examples:

Chat-GPT, LLaMA, Gemini, Mistral....

Decoder-only

causal self-attention





Encoder-only Family – Masked Language Model (MLM)

Mask: **Full attention** → all tokens can attend bidirectionally

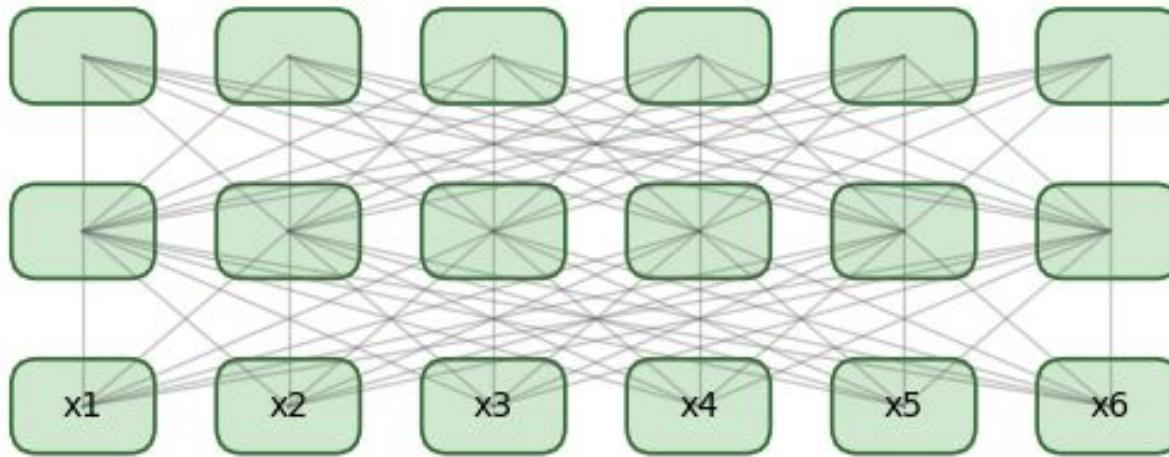
Objective: Predict **randomly masked tokens** using full context

Key: “Understands” but does **not** generate

Examples:

BERT, RoBERTa, SpanBERT, ALBERT.....

Encoder-only bi-directional self-attention





Large Language Models (LLMs)

LLMs are:

- Transformer-based models
- Trained on massive text corpora - O(B) tokens scale
- Learn general-purpose language understanding and generation abilities

[Scaling Law] Model parameters, data, and compute → emergent intelligence

Key Characteristics:

- Trained with self-supervised next-token prediction
- Exhibit few-shot / zero-shot capabilities
- Can be further adapted via fine-tuning to many downstream tasks



From Architecture → Paradigm ⚙

Old View: “Architecture Defines Intelligence”

- Researchers focused on building smarter structures (LSTM, Transformer)
- But even elegant designs failed without scale and rich data

Modern View: “**Data + Algorithm** Define Capability”

- The same Transformer backbone can become BERT, GPT, or T5, depending on the algorithmic design (PT, SFT and RL) and massive data



From Architecture → Paradigm ⚙️

Key Insight:

- Architecture shapes foundation, but data and algorithm shape behavior
- The training paradigm now matters more than structural novelty

Takeaway:

Modern breakthroughs (ChatGPT, Gemini) come not from new layers, but from **how we teach and scale the same architecture, with better data**

Analogy: Human Learning



Stage	Human	Model (LLM)
Pre-training	Learning from reading, observation, and experience about the world	Trained on massive text corpus - learns language, facts, and world knowledge
Fine-tuning	Getting job-specific/domain training (e.g., becoming an engineer etc.)	Task-specific SFT on labeled data - adapts to domains or strengthens reasoning
Reinforcement Learning	Socialization and feedback - learning norms, politeness, ethics from others	Preference tuning (RLxF) - adjusts to ANY preference signals

Flow: Foundation → Specialization → Alignment



LLM Pre-training



What Is Pre-training?



Train on a **massive unlabeled text corpus**

- Billions of words from books, and web pages - no manual labeling needed
- The model learns patterns of syntax, and world knowledge implicitly

Self-supervised objective:

- Does **not** need label
- Predict **masked (Encoder-only)** or **next tokens (Decoder-only)** from context
- The task creates its own supervision signal by **hiding or shifting tokens**



Pre-training Examples

Example 1 - Masked Language Modeling (BERT-style)

Input:

"The cat [MASK_1] on the [MASK_2]."

Target:

"The cat **sat** on the **mat**."

👉 The model must use bidirectional context to predict the masked token



Pre-training Examples

Example 2 - Next Token Prediction (GPT-style)

Input:

"The cat sat on the ?"

Target:

Next token = mat

👉 The model predicts the next word given all previous ones (causal connection)



What Is Pre-training?



Outcome:

The model learns a **general-purpose representation of language**

- capturing grammar, meaning, and relationships between entities

Result:

The pre-trained model becomes:

- a reusable foundation
- easily adapted to many downstream tasks (classification, summarization, etc.) through **fine-tuning**



Why Self-Supervised? 🎉

Training signal comes directly from the data itself

Every word, punctuation, or phrase provides a new learning opportunity

This makes the world knowledge effectively one massive training corpus

Scales effortlessly:

Since labeling is automatic, models can train on trillions of tokens across diverse domains
(books, code, social media etc)



Pre-training Data



High-Quality Curated Web Data

- StackExchange / Reddit / Hacker News / ArXiv / PubMed - filtered for long-form reasoning, technical correctness, and linguistic quality
- C4 (Colossal Clean Crawled Corpus) - curated Common Crawl derivative used in T5
- RefinedWeb / Dolma / RedPajama / Falcon RefinedWeb - open-source large-scale cleaned web corpora



Pre-training Data



Instructional / Human-Curated Data

- Wikipedia discussions, StackOverflow Q&A, instructional forums — naturally structured in “question → answer” form
- Open-sourced QA datasets: Natural Questions, SQuAD, GSM8K, MMLU-style evaluation corpora

Educational and Academic Texts

- Textbooks, lecture notes, academic papers (via S2ORC or ArXiv) - help models internalize reasoning and math syntax
- Project Gutenberg + scientific book scans - classic long-form, well-structured writing



Pre-training Data



Code and Technical Content

- GitHub, StackOverflow, Jupyter Notebooks, Competitive Programming datasets

Dialogue & Conversational Corpora

- OpenSubtitles, MultiWOZ, OpenOrca - to enhance conversational fluency

Multilingual and Multimodal Text

- Multilingual datasets: Wikipedia in 100+ languages - encourage cross-lingual generalization
- Multimodal text: image captions, audio transcripts, or paired HTML/context



Pre-training Data



Synthetic & Augmented Data (increasingly common)

Self-generated text from earlier LLMs (self-play, bootstrapping, or distillation)

Data augmentation via back-translation, paraphrasing, or chain-of-thought synthesis

Steps: Data cleaning → Deduplication → Quality filtering

Goals: Diversity, quality, and safety → Teach LLMs the basic knowledge



Pre-training Data Cleaning Pipeline



Deduplication

- Remove exact and near-duplicate documents using MinHash or SimHash

Content Filtering

- Exclude boilerplate (HTML templates, navigation bars, ads)

Toxicity & Safety Filtering

- Use classifiers to detect hate, harassment, adult, or violent content
- Combine rule-based filters with model-based toxicity scoring



Pre-training Data Cleaning Pipeline



Language & Format Detection

- Identify language via FastText or Compact Language Detector (CLD3)
- Keep target languages; drop mixed or unrecognized ones

Quality Scoring & Sampling

- Train classifiers or use perplexity thresholds to rank document quality.
- Sample more from high-quality sources (books, Wikipedia) than noisy web

Normalization & Finalization

- Lowercase normalization, punctuation cleanup, HTML stripping
- Mask or remove personally identifiable information (PII)



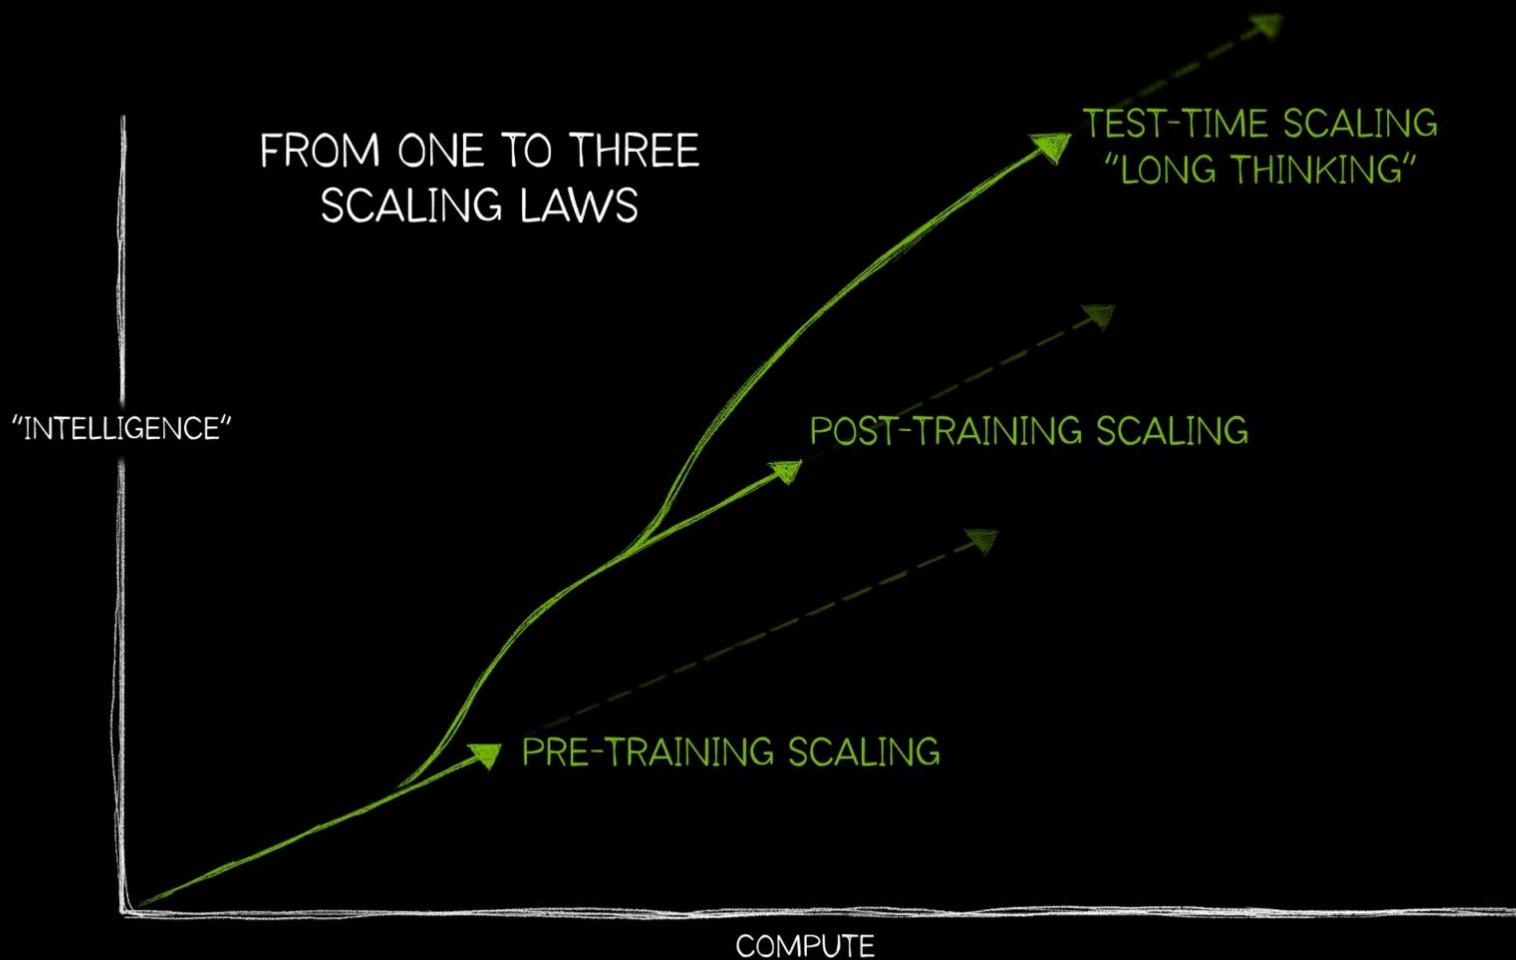
Scaling Laws



Model performance improves predictably as model size, data, and compute increase - following a power-law trend

Patterns observed:

- More data → smoother loss curves, better generalization
- Larger models → capture richer structure, longer reasoning chains
- Emergent abilities appear only after certain scale thresholds
- But returns diminish - doubling data \neq doubling capability





Why Next-Token Prediction Matters 😐

The next token prediction task is extremely simple

Yet applied at **massive scale** the model **implicitly** learns many tasks

Grammar: “In my free time, I like to {run, banana}” → predicting “run”

World knowledge: “The capital of Japan is {Copenhagen, Tokyo}” → predicting “Tokyo”

Sentiment: “I was laughing the entire time, the movie was {good, bad}” → predicting “good”

Reasoning & Logic: “If it rains, the ground will be {wet, dry}.” → wet

Math & Symbolic Reasoning: “ $2 + 3 = \{5, 6\}$ ” → 5



What the Model Learns



Syntax / Semantics / Factual knowledge / Basic reasoning capabilities /

Each next-token prediction is like solving a **micro-task**

- fill in a blank, finish a phrase, complete a thought

Across **billions of sentences**, the model encounters millions of **implicit tasks** -
translation, analogy, classification, reasoning, summarization

The model does NOT memorize text; it models **probability** of the next token

- learning which **continuations** are plausible, logical, and coherent



Emergent capability



At sufficient scale, next-token prediction yields general intelligence-like behavior

- models can perform tasks they were never explicitly trained for, by transfer what they have learned from languages

This phenomenon is often summarized as:

- Pre-training teaches basic knowledge
- Supervised Fine-tuning enhance reasoning capabilities
- Reinforcement Learning aligns it

General Artificial Intelligence Step 0



Transformer + next-token objective + large data

→ General-purpose language model

Emergent abilities appear once enough **scale + data + training**



LLM Supervised Fine-tuning



From Pre-training to Fine-tuning



Pre-training: Learns general world knowledge from massive unlabeled text

- Models know how to speak but not what to do

Fine-tuning: Teaches behavior and task goals using labeled examples

- Bridge: Moves from “foundation knowledge” → “actionable skills”



Supervised Fine-tuning (SFT)

Train a pre-trained LLM on **(prompt, response)** pairs for a target task/behavior

Goal: Align the model to follow instructions and produce useful, reliable outputs

Setup:

- Freeze nothing (Full SFT) or lightly freeze ([PET](#))
- Minimize **token-level cross-entropy** on target responses

Data: High-quality, diverse instructions with clear, verifiable targets



What Does the SFT Loss Do? 🤔

Cross-entropy loss = negative log probability of the correct next token

$$L = - \sum \log P(\text{correct_token} | \text{prompt}, \text{previous_tokens})$$

Each token in the reference answer contributes to the loss

If the model assigns low probability to a correct token

- higher loss → stronger learning signal



Standard SFT Paradigm

Example:

Prompt: “Summarize the paragraph.”

Response: “Climate change is harming coral reefs.”

Model learns to:

- Follow instructions
- Generate fluent, polite, and relevant text

Limitation: Learns what to say, **not** how to think



Why Modern LLMs Are “Thinking Models”

From Pattern Matching → Structured Reasoning

- Early LLMs mimicked language fluency
- Newer models learn to analyze, plan, and reason across multiple steps

Training Evolution

- Pre-training on vast, diverse data teaches latent world models
- SFT and RL explicitly train reasoning traces (“think step by step”)

Behavior Shift

- Models now explain reasons before answering
- Reasoning improves consistency, interpretability, and self-correction



Thinking SFT is Step 1

Standard SFT = aligns immediate behavior (instruction-following)

Thinking SFT = aligns reasoning (how the model reaches conclusions)

Without reasoning signals, models learn to “answer fast” instead of “think slow”

Goal: expose the model to step-by-step problem solving, not just end results



We are teaching the model to show its work - not just the answer



What Are “Thinking Tasks”?

Tasks where intermediate reasoning is essential, not optional

Example 1 - Math / Logic

Prompt: “What is 27×14 ?”

Reasoning: “ $27 \times 10 = 270$, $27 \times 4 = 108$, $270 + 108 = 378$.”

Answer: “378.”

Example 2 - Everyday Reasoning

Prompt: “Why do we wear jackets in winter?”

Reasoning: “Because jackets trap body heat, keep us warm when the air is cold.”

Answer: “To stay warm.”



Example 1 – Plain SFT 🧱

Prompt: “Solve: $12 + 7 = ?$ ”

Response: “19.”

Loss is computed only on the final tokens “1”, “9”, “.”

- The model learns to produce correct answers and follow format
- It does not learn how to reason



Example 2 – Reasoning-augmented SFT

Prompt: “Solve: $12 + 7 = ?$ ”

Response: “Let’s think step by step. $12 + 7 = 19$. Answer: 19.”

Now the loss covers every token, including

“Let’s”, “add”, “step”, “by”, “step”, “12”, “+”, “7”, “=”, “19”...

- The model learns that generating intermediate reasoning steps reduces the overall loss
- The behavior “explain before answering” becomes part of training



Example 3 – Weighting by Reasoning Quality

Prompt: “Solve: $12 + 7 = ?$ ”

Sample Text

- | | |
|---|---------------------------|
| 1 | “I think $12 + 10 = 22$ ” |
| 2 | “ $12 + 7$ shall be 19” |

Correct?

- | |
|--|
| |
| |

Loss can be weighted across samples:

$$L = w_1 \times CE(\text{sample}_1) + w_2 \times CE(\text{sample}_2)$$

Only correct or verified samples receive higher weight

→ The model learns which reasoning paths are trustworthy



Where to Source Thinking Data ⚙

Math / Logic: GSM8K, MATH, AQUA-RAT, SVAMP

Reasoning QA: StrategyQA, HotpotQA, OpenBookQA

Step-by-step: CoT (Chain-of-Thought) traces from Flan, CotHub, OpenOrca

Code Reasoning: CodeContests, LeetCode-HF, DeepMind MathCode datasets

Human-curated: Collect instructor-written rationales or worked solutions

Synthetic Generation: Use an existing LLM to produce reasoning traces

("Let's think step by step..." → label + verify)

Filter by correctness or **self-consistency**



How to Curate and Format



Keep structure consistent

Prompt: [problem]

Reasoning: [chain of thought]

Answer: [final result]

Keep reasoning concise

2–6 reasoning steps preferred

Each step should be verifiable (not vague “I think maybe...”)



How to Curate and Format



Keep correctness

Human or model verification for every rationale

Discard hallucinated or logically invalid traces

Keep diversity

Mix numeric, symbolic, commonsense, procedural reasoning

Include both short and long-context tasks



Beyond Single-task SFT



Multi-task fine-tuning: Train on many instruction types simultaneously

Mixture of formats: QA + translation + summarization + more → boosts skills

Instruction format standardization: Convert all tasks to prompt-response pairs

Results: Better zero-shot transfer and alignment readiness



Industry trends



Looking Ahead

The pre-training + fine-tuning paradigm keeps evolving rapidly

New research improves **efficiency**, **scalability**, and **capability**

Understanding these trends helps you stay ahead of the LLM curve

Focus today:

- Efficiency
- Multimodality
- Domain adaptation



Parameter-efficient Fine-tuning (PEFT)

Challenge: Full fine-tuning = billions of parameters → expensive & slow

Idea: Freeze the base model; **only train small additional modules**

Goal: Reduce compute + memory footprint while keeping performance

Result: Same model reused across many downstream tasks efficiently

LoRA: Low-Rank Adaptation



Key idea: Instead of updating full weight matrix W , learn 2 small matrices A & B s.t.

$$\Delta W = A \times B^T \text{ (low-rank update)}$$

Benefits:

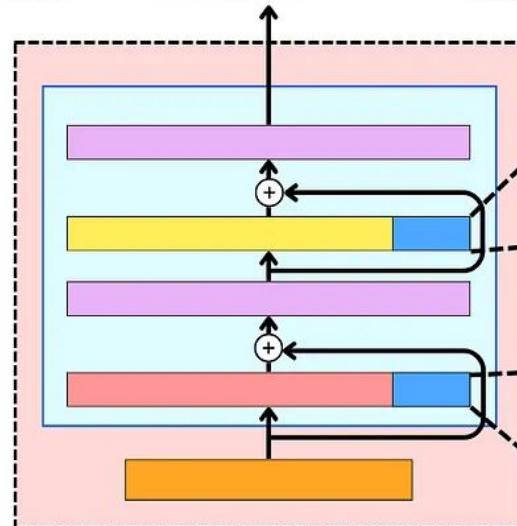
- ~1–2% trainable parameters
- Easy to merge/unmerge with base model
- Great for **multi-task learning** and **low-data** scenarios

$$\theta_F = \boxed{\theta_T} + \boxed{\Delta W} \simeq \theta_T + \boxed{B^T A}$$

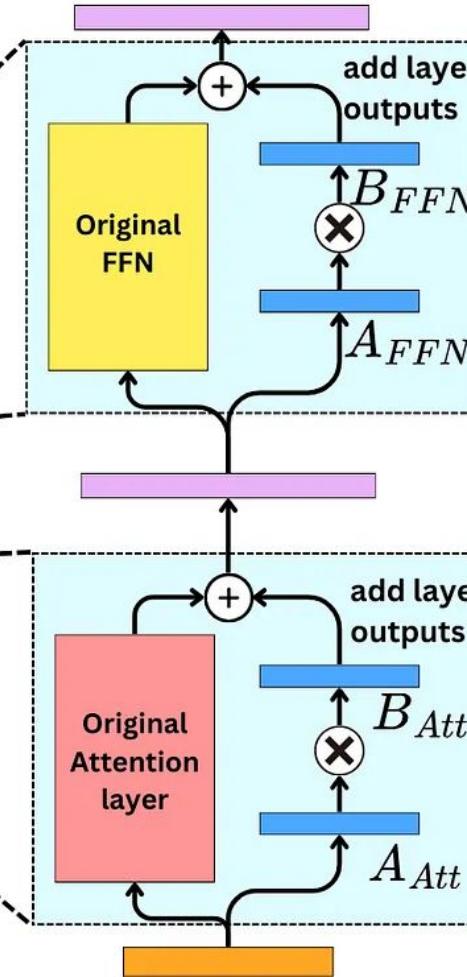
The trained
model

The fine-
tuning data

Low-Rank
approximation

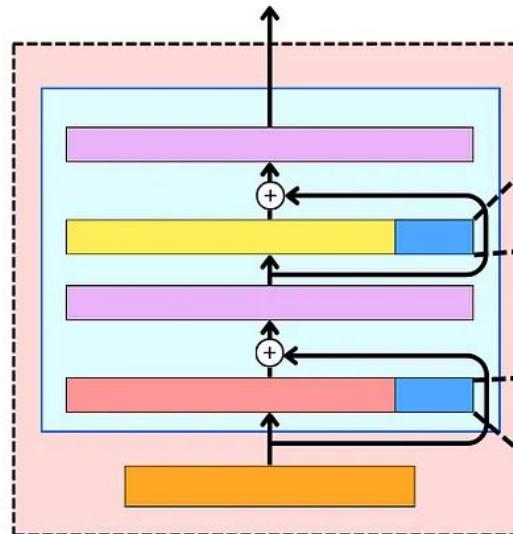


Forward pass

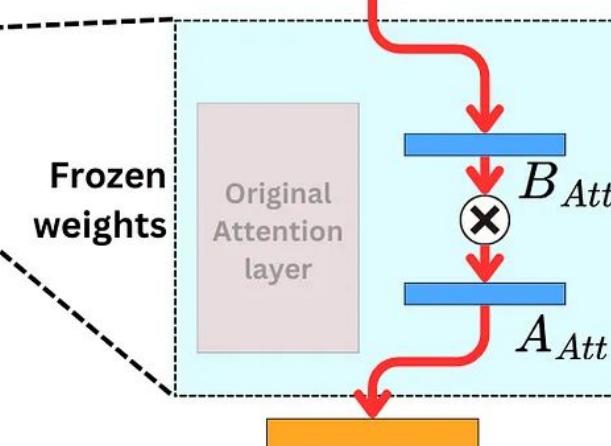
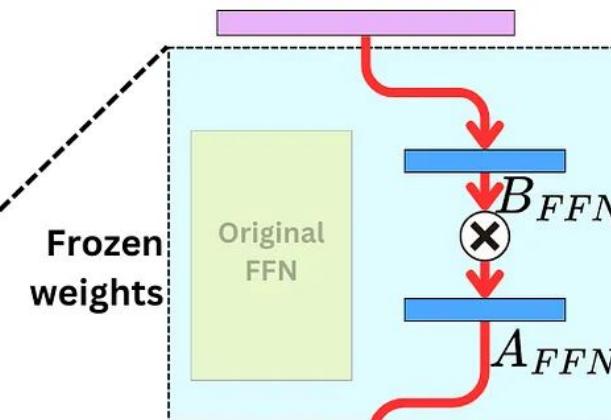


$$A_t \leftarrow A_{t-1} - \alpha \nabla_A \mathcal{L}|_{A=A_{t-1}}$$

$$B_t \leftarrow B_{t-1} - \alpha \nabla_B \mathcal{L}|_{B=B_{t-1}}$$



Augment layers with
LoRA Adaptors



Backward pass



QLoRA: Quantized LoRA ⚡

Problem: Even LoRA can be memory-heavy for >30B models

Quantize frozen/trainable weights to 4-bit precision → only train LoRA layers

Outcome: Fine-tune 65B models on a single 48GB GPU 😺

Used in: almost all open-source chat models



Multimodal Pre-training



Text + Image → CLIP

Contrastive objective aligns visual & textual embeddings

Text + Vision + Audio → Flamingo, Gemini, ChatGPT

Cross-modal attention for unified reasoning

Why It Matters:

Language models are expanding beyond text - “seeing” and “hearing”

Ongoing Effort: **Unified foundation models across modalities**



Domain Adaptation



Tailor models to specialized domains (e.g., law, finance, biomedicine)

Approach: Further tuning (SFT/RL) on domain-specific corpora

Examples:

- CodeLlama → trained on github data to deal with coding problems
- BloombergGPT → trained on financial news, and market data to handle finance-specific tasks (analysis, risk, compliance)
- Med-PaLM / Meditron → fine-tuned on medical QA datasets for clinical reasoning

Impact: Improves factual accuracy and domain reasoning



Summary

Pre-training: foundation of general intelligence

SFT: teaches forcing behavior and step 1 of thinking LLMs

Emerging trends: efficiency, multimodality, domain focus

Mindset shift: Foundation models are not endpoints - but platforms

Next Lecture: Reinforcement Learning for “AGI” 🤔